

Devops Assignment 2 - Yousef Amirghofran

(Briefbot)

In the first part of the assignment, I built a great project called Briefbot that allowed users to submit links via the web page or a browser extension and an AI-enabled workflow would then extract information from the link, summarize it, and provide a podcast about the item using external API services.

In this assignment, I improved the code quality and pursued DevOps practices to deploy the project on Azure.

The report will focus on explaining each part of this process in better detail with examples.

Refactoring Code Smells

After learning about SOLID principles and code smells, I reviewed the code in my project to see whether I could identify any code smells and improve the quality of the code. And I found a lot of places where I could improve the code:

Bloated Handler (SRP Violation)

```
type Handler struct {
    userService      services.UserService
    itemService     services.ItemService
    digestService   services.DigestService
    podcastService  services.PodcastService
    sseManager      *services.SSEManager
}
```

I had created a handler struct in my project that would receive each service (userService, itemService, digestService, podcastService, and sseManager) as an input. This made the handler depend on 5 different services and handle the routing for multiple domains (users, items, podcasts, and digests) which violated the Single Responsibility Principle as the struct had too many reasons to change. Any change in any of those services could potentially propagate the change to the handler struct.

To fix this problem, I refactored the handler struct and broke it down into more focused structs where each handler dealt with its own domain. This made the code much easier to maintain and scale later on.

```
type UserHandler struct {
    userService services.UserService
}

type ItemHandler struct {
    itemService services.ItemService
    sseManager *services.SSEManager
}

type PodcastHandler struct {
    podcastService services.PodcastService
    sseManager     *services.SSEManager
}

type DigestHandler struct {
    digestService services.DigestService
}
```

Hardcoded Worker Service (OCP Violation)

I had written a `processBatch` function in the worker service for the workers to process the jobs of creating a podcast for the items. However, I had hardcoded the processing of the podcasts from the podcasts service, tightly coupling the service that created the jobs with the worker service. This made it so that any time I wanted to add a new type of job or service to be processed, I had to edit the `processBatch` function which broke the Open-Closed principle where the function should be open to extension but closed for modification

```
func (s *workerService) processBatch() error {
    // Process items first
    if err := s.processItemBatch(); err != nil {
        return fmt.Errorf("failed to process item batch: %w", err)
    }

    // Process podcasts if enabled
    if s.enablePodcasts && s.podcastService != nil {
        if err := s.processPodcastBatch(); err != nil {
            return fmt.Errorf("failed to process podcast batch: %w", err)
        }
    }
}
```

```
    return nil
}
```

To fix this, I decided to use the strategy pattern in the `processBatch` function and used Dependency Injection to pass a list of the processors to the worker service, decoupling the services that create the jobs from the worker service. Now, I don't have to edit the function when I add a processor, it automatically handles it when I pass in the processor when initiating the service.

```
type JobProcessor interface {
    Process(ctx context.Context, batchSize int32) error
}

type workerService struct {
    jobQueueService JobQueueService
    processors      []JobProcessor
    // ... other fields
}

func (s *workerService) processBatch() error {
    for _, processor := range s.processors {
        if err := processor.Process(s.ctx, s.batchSize); err != nil {
            return err
        }
    }
    return nil
}
```

These two examples serve to represent the larger number of refactorings of similar nature that I did to improve the code quality. Overall, these fixes significantly improved the quality of my code and made it more maintainable and extensible for the future.

Testing

I followed the testing conventions in Go in this project, writing tests in `_test.go` files and using the rich go testing library.

I started with writing unit tests that verified functionality of small isolated functions. These tests shape the majority of the tests because they are easier and faster to write and run (which is important later when we run them in the CI/CD pipeline with every pull request).

At this stage, I also wrote mock tests. For that, I created a centralized test utilities package where I defined the fixtures, mock interfaces, and helper test functions for writing mock tests. For example, the `MockQuerier` interface mocked real database operations and the `TestDataBuilder` function used the Factory method to create the test data for certain mock tests. Since my project

uses a lot of complex external APIs, I was limited in the number of services I could mock. That's why my code coverage wasn't as high as others at ~70%.

However, I also wrote integration tests to test the actual behavior of some of the services I was using such as Cloudflare R2 for object storage and FFmpeg for audio stitching. There were still some limitations to test the Fal service since it's a very complicated API.

Over the different parts of the project (metrics, handlers, middleware, and services), I reached a test coverage of ~85% with all tests passing. Test reports are included in the repo at `backend/test-results.txt` and `coverage.html`.

Containerization

To prepare for the deployment of this project, I containerized the project using Docker. First, I created multi-stage Dockerfile builds for the backend (Go API) and frontend (React Tanstack Start) components of the project to enable creating docker images for each of them. Then, I created a `docker-compose.yml` file at the root of the project to orchestrate all the interconnected services including the backend, frontend, PostgreSQL database, and the monitoring infrastructure (Prometheus and Grafana images and configurations).

This not only allowed me to build and run the project much faster with a single `docker-compose up` script, but it enabled me to configure the deployment pipeline as well and automate the deployment as well.

Monitoring

After a bit of research, I realized that the most common stack for monitoring is using Prometheus and Grafana where Prometheus acts as a time-series database that scrapes the project for metrics such as latency, error rate, and health status and Grafana is an easy way to declaratively create the graphical dashboards using YAML files to display the metrics in a visual dashboard.

To do this, I first added a simple `/health` endpoint to my backend API that returned success if the backend was up. Then, I created a `metrics.go` package to instrument my backend functions with metrics that would be exposed through the new `/metrics` endpoint. Additionally, I used the Prometheus Go client library to create a middleware to automatically log request metrics without changing the business logic. This made it so that Prometheus could scrape the metrics from my app. Then, I configured the Prometheus settings in `monitoring/prometheus.yml` to scrape metrics in 15-second intervals. Here are the main metrics I collected with Prometheus:

- HTTP: request counts, latency, request/response sizes, active requests

- Job Queue: num. enqueued tasks, num. processing tasks, num. completed tasks, num. failed tasks
- Workers: num. active workers, num. jobs processed per worker, worker error rates, batch durations
- Podcasts: counts, durations, audio requests, failure rate
- SSE: num. active connections, num. events sent by type

Under `monitoring/grafana/`, I set up the Grafana dashboard structure and determined a number of panels with PromQL queries to aggregate and visualize the data from Prometheus.

Finally, in the docker-compose file, I set up both the Prometheus and Grafana services with persistent volumes (`prometheus_data` and `grafana_data`) and configured the ports across the `briefbot-network` bridge network.

CI/CD Pipeline

I used Github Actions as my CI/CD pipeline solution because my code already resides in Github and the most natural way to perform Continuous Integration, and later Deployment, is in the same platform.

There are multiple stages in the CI/CD pipeline. First, it runs lint and formatting checks on the code. This is standard practice to make sure all code follows the projects' conventions and allows developers to focus on the logic of the code, rather than noisy formatting inconsistencies. After that, it builds them to verify they compile.

Next, since I use `SQLc` to auto-generate Go functions from SQL files, I use the `verify` function from `SQLc` to verify my queries are in sync with my database schema. This catches potential database migration errors before they can cause any problems so they can be fixed in the Pull Request itself.

Then, the pipeline runs the test suite. It both verifies that all tests pass and the overall test coverage is above 70%. If these conditions aren't true, it fails and the developer can see how to fix the issues. This also helps prevent pushing errors to production.

The previous stages run on any pull request. On any push to the `main` branch, after all the previous stages succeed, the pipeline builds the docker containers, as discussed previously, and publishes them to Docker Hub (the authentication tokens for Docker Hub are configured in Github Secrets). This allows our deployment provider to pull these containers every time they are updated and deploy them.

Deployment

Finally, everything is ready to deploy the project. I thought the best Azure service I could use was the App Service since it works with containers and natively integrates with monitoring products such as Azure Insights.

I created two Azure App Services, one for the frontend component of my app, and one for the backend. In both cases, I configured the container registry as Docker Hub and provided the name of the containers to the corresponding services (backend and frontend) with the `latest` tag. Every time my CI/CD pipeline publishes a new container, the App Services pull them and deploy them.

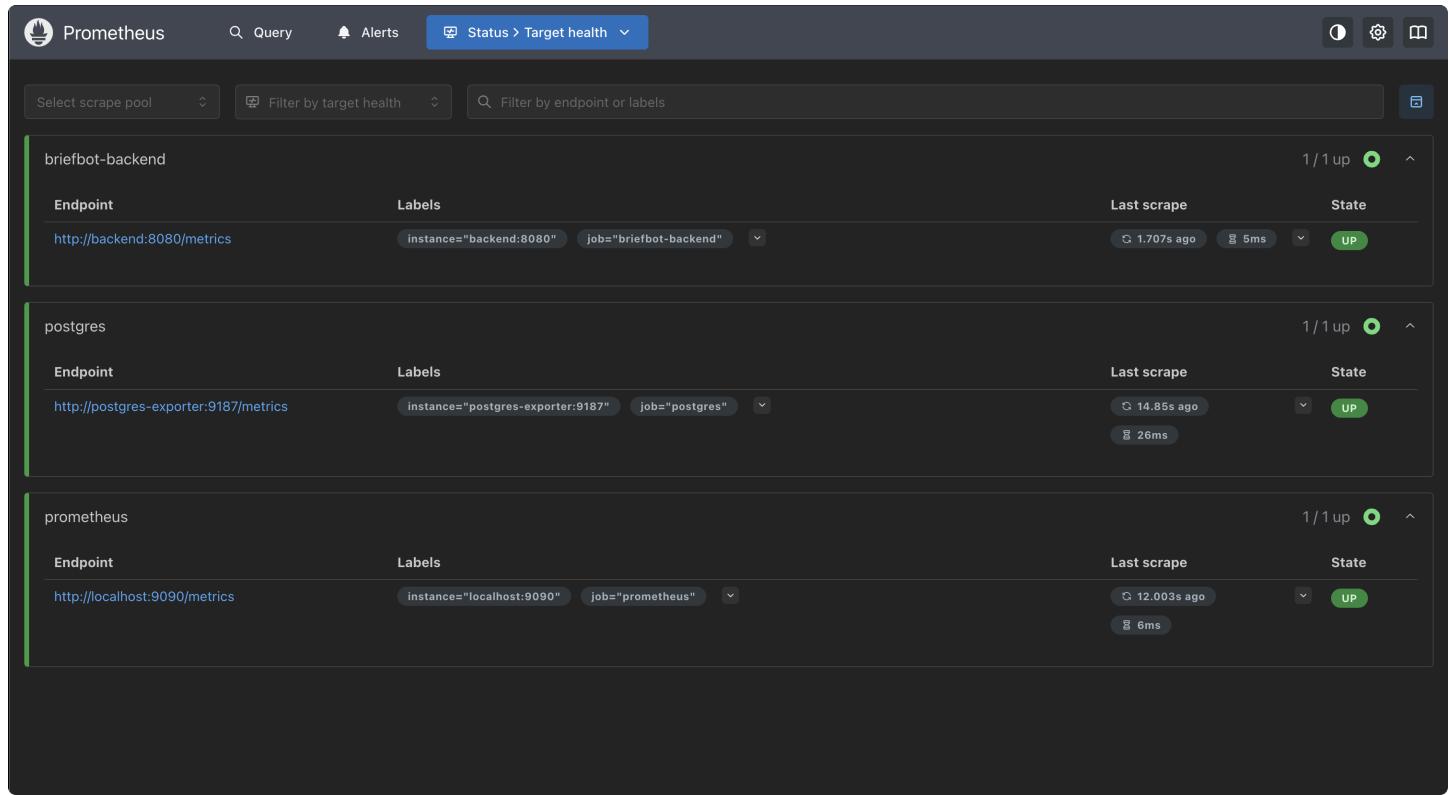
Finally, I created a Flexible PostgreSQL instance in Azure and use the URL for it as the `DATABASE_URL` environment variable in my backend service to connect the backend to the database.

Conclusion

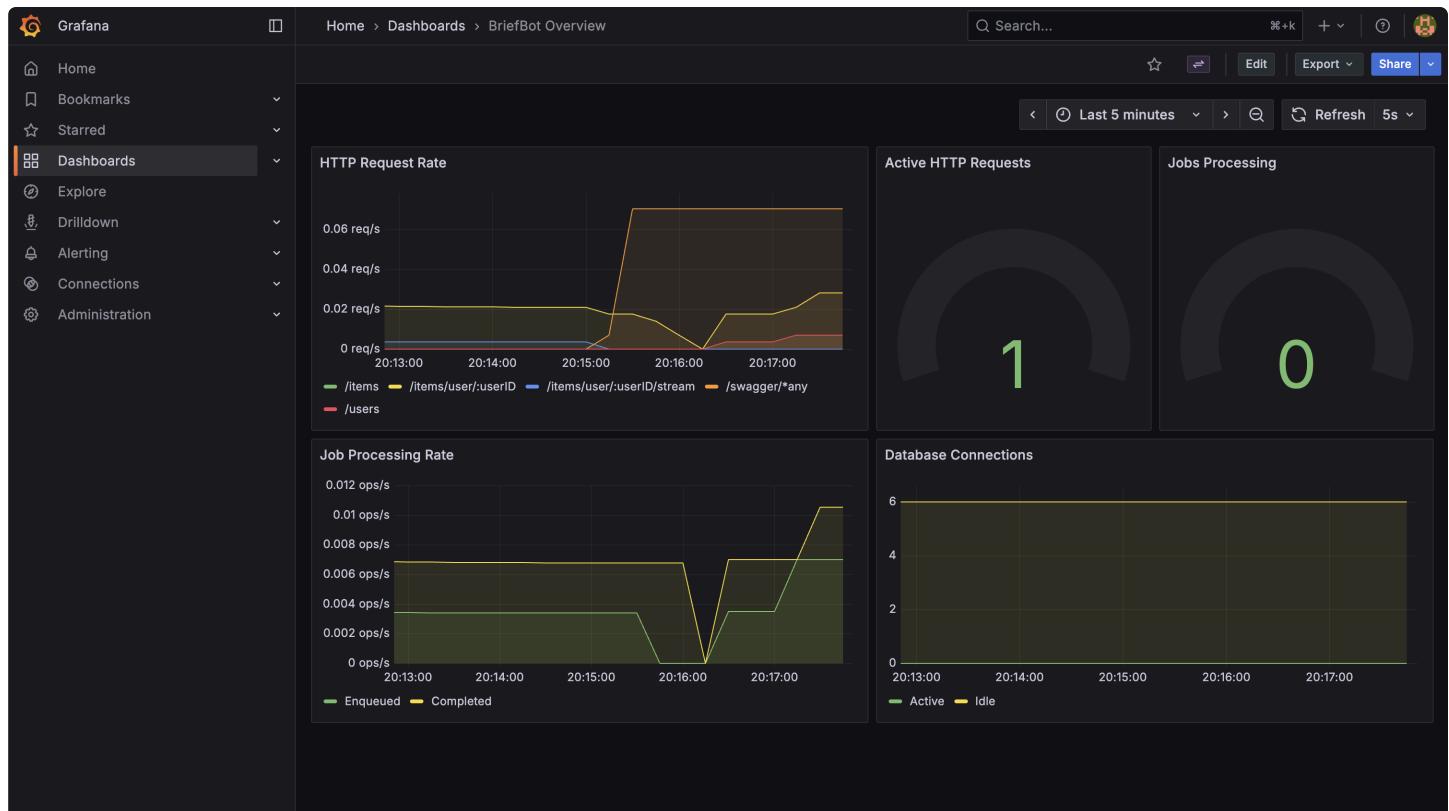
In this assignment, I reviewed the project to refactor certain code smells to improve the quality of the code and make it compliant with SOLID principles, added unit and integration tests to ensure the robustness and reliability of the project, containerized the different components for easier deployment, added monitoring using Prometheus and Grafana, and created a CI/CD pipeline to automate deployment to the Azure cloud.

Appendix

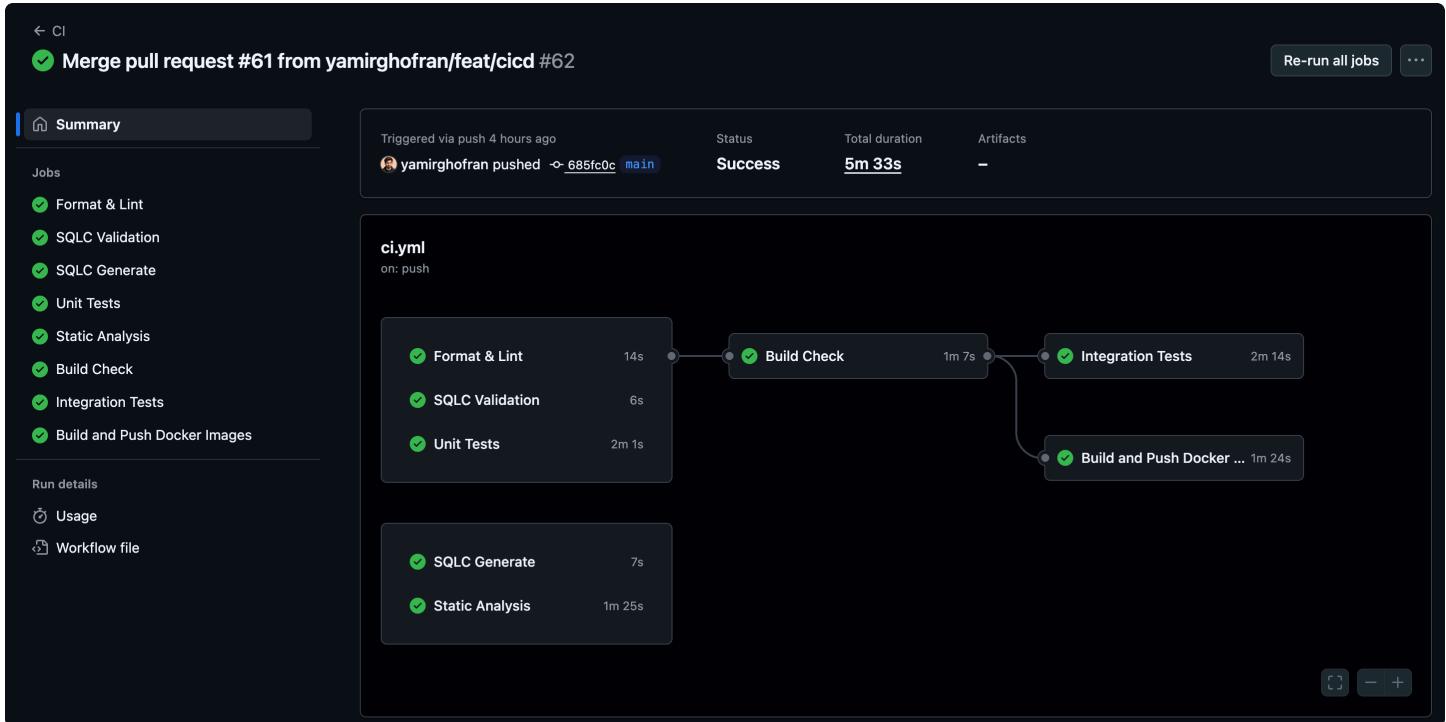
Prometheus Dashboard



Grafana Dashboard



Github Actions CI/CD



yamirghofran/briefbot-frontend

Last pushed about 23 hours ago • Repository size: 50.6 MB • ⭐0 • ⌂47

[Add a description](#) [Add a category](#) [General](#) [Tags](#) [Image Management](#) BETA [Collaborators](#) [Webhooks](#) [Settings](#) **Tags** DOCKER SCOUT INACTIVE
[Activate](#)

This repository contains 3 tag(s).

Tag	OS	Type	Pulled	Pushed
main-685fc0c		Image	less than 1 day	about 23 hours
latest		Image	less than 1 day	about 23 hours
main		Image	less than 1 day	about 23 hours

[See all](#)Build with
Docker Build Cloud

Accelerate image build times with access to cloud-based builders and shared cache.

Docker Build Cloud executes builds on optimally-dimensioned cloud infrastructure with dedicated per-organization isolation.

Get faster builds through shared caching across your team, native multi-platform support, and encrypted data transfer - all without managing infrastructure.

[Go to Docker Build Cloud](#)

Using 0 of 1 private repositories.

[Repository overview](#) INCOMPLETE

Repositories / briefbot-backend / General

yamirghofran/briefbot-backend

Last pushed about 23 hours ago • Repository size: 95.3 MB • ⭐0 • ⌂78

[Add a description](#) [Add a category](#) [General](#) [Tags](#) [Image Management](#) BETA [Collaborators](#) [Webhooks](#) [Settings](#) **Tags** DOCKER SCOUT INACTIVE
[Activate](#)

This repository contains 4 tag(s).

Tag	OS	Type	Pulled	Pushed
main-685fc0c		Image	less than 1 day	about 23 hours
latest		Image	less than 1 day	about 23 hours
main		Image	less than 1 day	about 23 hours
main-6abddd31		Image	less than 1 day	about 23 hours

[See all](#)Build with
Docker Build Cloud

Accelerate image build times with access to cloud-based builders and shared cache.

Docker Build Cloud executes builds on optimally-dimensioned cloud infrastructure with dedicated per-organization isolation.

Get faster builds through shared caching across your team, native multi-platform support, and encrypted data transfer - all without managing infrastructure.

[Go to Docker Build Cloud](#) [Repository overview](#) INCOMPLETE**Azure App Services**

The screenshot shows the Azure portal interface for the 'briefbot-backend' web application. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Microsoft Defender for Cloud, Events (preview), Resource visualizer, Deployment (Deployment slots, Deployment Center), and Settings (Environment variables, Configuration (preview), Instances, Authentication, Identity, Backups, Custom domains, Certificates, Network). The main content area displays the 'Essentials' section with details like Resource group (BCSAI2025-DEVOPS-STUDENTS-A), Status (Running), Location (West Europe), Subscription (Azure Simple IF Instituto de Empresa, S.L.), Subscription ID (e0b9cada-61bc-4b5a-bd7a-52c606726b3b), and Tags (Add tags). Below this is a 'Properties' tab with sections for Web app, Domains, and Hosting, each with specific configuration details. A 'Notifications' tab shows 2 notifications. The top right has a 'JSON View' link. The bottom navigation bar includes Copilot, Search resources, services, and docs (G-), and the URL yamirghofran.eu2023@... IE INSTITUTO DE EMPRESA, S.L.

The screenshot shows the Microsoft Azure portal interface. The top navigation bar includes the Microsoft Azure logo, a search bar, and various global settings and notifications. The main title is "briefbot-backend | Environment variables" under the "Web App" category. On the left, a sidebar lists navigation options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Microsoft Defender for Cloud, Events (preview), Resource visualizer, Deployment (with slots and Center), and Settings (with Environment variables selected). The main content area displays a table of environment variables with columns for name, value visibility, service type, and edit icon. The variables listed include FAL_API_KEY, GROQ_API_KEY, R2_ACCESS_KEY_ID, R2_ACCOUNT_ID, R2_BUCKET_NAME, R2_PUBLIC_HOST, R2_SECRET_ACCESS_KEY, SES_FROM_EMAIL, SES_FROM_NAME, SES_REPLY_TO_EMAIL, TELEGRAM_BOT_TOKEN, WEBSITES_ENABLE_APP_SERVICE_STORAGE, and XDT_MicrosoftApplicationInsights_Mode, all categorized as App Service.

Name	Show value	Type	Action
FAL_API_KEY	Show value	App Service	
GROQ_API_KEY	Show value	App Service	
R2_ACCESS_KEY_ID	Show value	App Service	
R2_ACCOUNT_ID	Show value	App Service	
R2_BUCKET_NAME	Show value	App Service	
R2_PUBLIC_HOST	Show value	App Service	
R2_SECRET_ACCESS_KEY	Show value	App Service	
SES_FROM_EMAIL	Show value	App Service	
SES_FROM_NAME	Show value	App Service	
SES_REPLY_TO_EMAIL	Show value	App Service	
TELEGRAM_BOT_TOKEN	Show value	App Service	
WEBSITES_ENABLE_APP_SERVICE_STORAGE	Show value	App Service	
XDT_MicrosoftApplicationInsights_Mode	Show value	App Service	

Microsoft Azure

Home > briefbot-frontend Web App

briefbot-frontend Web App

Search ...

Browse Stop Swap Restart Delete Refresh Download publish profile Reset publish profile Share to mobile Send us your feedback

You have multiple notifications. Switch to Notifications tab to get more details.

Overview

- Activity log
- Access control (IAM)
- Tags
- Diagnose and solve problems
- Microsoft Defender for Cloud
- Events (preview)
- Resource visualizer
- Deployment
- Deployment slots
- Deployment Center
- Settings
- Environment variables
- Configuration (preview)
- Instances
- Authentication
- Identity
- Backups
- Custom domains
- Certificates
- Networking

Add or remove favorites by pressing Cmd+Shift+F

Essentials

Resource group (move) : BCSAI2025-DEVOPS-STUDENTS-A	Default domain : briefbot-frontend.azurewebsites.net
Status : Running	App Service Plan : ASP-BCSAI2025DEVOPSTUDENT1A-ab55 (B1: 1)
Location (move) : West Europe	Operating System : Linux
Subscription (move) : Azure Simple IE Instituto de Empresa, S.L.	Health Check : Not Configured
Subscription ID : e0b9cada-61bc-4b5a-bd7a-52c606726b3b	

Tags (edit) : Add tags

Properties Monitoring Logs Capabilities Notifications (2) Recommendations

Web app

Name : briefbot-frontend	Deployment Center
Publishing model : Container	Deployment logs
Container Image : index.docker.io/yamirghofran/briefbot-frontend:latest	View logs
Runtime status : Healthy	

Domains

Default domain : briefbot-frontend.azurewebsites.net	Virtual IP address : 20.50.2.94
Custom domain : Add custom domain	Outbound IP addresses : 135.236.89.194, 108.141.27.65, 20.67.59.47, 9.163.255.242, 132.220.19.8, 108.142.59.236, 128.251.211.129, 132.220.25.182, 98.64.12.115, 132.164.19.45, 57.153.161.115, 9.163.79.192, 20.50.2.94

Hosting

Plan Type : App Service plan	Additional Outbound IP addresses : 135.236.89.194, 108.141.27.65, 20.67.59.47, 9.163.255.242, 132.220.19.8, 108.142.59.236, 128.251.211.129, 132.220.25.182, 98.64.12.115, 132.164.19.45, 57.153.161.115, 9.163.79.192, 20.50.2.94
------------------------------	--

Deployment Center

Application Insights

Networking

JSON View

Azure PostgreSQL

Microsoft Azure

Home > briefbot-db Azure Database for PostgreSQL flexible server

briefbot-db Azure Database for PostgreSQL flexible server

Search ...

Connect from VS Code Delete Reset password Restore Restart Upgrade Stop Refresh CLI / PS Feedback

Diagnose issues with this resource Generate Bicep template for this resource Check health status of this resource

Overview

- Activity log
- Access control (IAM)
- Tags
- Diagnose and solve problems
- Resource visualizer
- Migration
- Fabric mirroring
- Settings
- Power Platform
- Security
- Intelligent Performance
- Monitoring
- Automation
- Help

Add or remove favorites by pressing Cmd+Shift+F

Essentials

Subscription (move) : Azure Simple IE Instituto de Empresa, S.L.	Endpoint : briefbot-db.postgres.database.azure.com
Subscription ID : e0b9cada-61bc-4b5a-bd7a-52c606726b3b	Virtual endpoint : Not enabled
Resource group (move) : BCSAI2025-DEVOPS-STUDENTS-A	Administrator login : briefbot
Server name : briefbot-db	Configuration : General Purpose, D4ds v5, 4 vCores, 16 GiB RAM, 128 GiB storage
Location : West Europe	PostgreSQL version : 17.6
Status : Ready	Availability zone : 1
Created on : 2025-11-22 15:35:43.8177316 UTC	High availability : Not enabled

Tags (edit) : Add tags

Getting started Properties Recommendations (3) Monitoring Tutorials

Start your project

Connect to your database for the first time with a few simple steps.

Allow access Configure network access to your PostgreSQL database Configure Networking

Connect Learn how to connect with the application driver of your preference: .NET, PostgreSQL, Python, JDBC, PHP, Node.js, Ruby, PHP, C++ View connection strings

Samples Setup a sample database schema to get started Samples in GitHub

Mirroring in Microsoft Fabric Mirror your databases in Fabric, where you have access to all your data.

JSON View